

Offload the Coding Subscription

On-premises GPU serving for OpenCode on AMD Radeon™ AI PRO and Instinct™

The subscription is a meter

Commercial plans sell a seat, then bill the agent loop again once the included pool is gone. Enterprise often sells the seat as access only.


Offering, annual list price	Seat / month	3-year seat	After the included pool
GitHub Copilot Business	\$19	\$684	\$0.01 per AI credit
Cursor Teams, standard	\$40	\$1,440	API price, plus \$0.25 / M tokens
Claude Pro; Max from \$100	\$20; \$200	\$720; \$7,200	API rates
Claude Team, standard; premium	\$20; \$100	\$720; \$3,600	purchased extra capacity
Claude Enterprise	\$20	\$720	every token, at API rates
ChatGPT Business, standard; premium	\$20; \$100	\$720; \$3,600	workspace credits
ChatGPT Enterprise	custom	quote	credits or tokens

Claude Sonnet 5 list rate, the meter on Claude Enterprise, is **\$2 / \$10** per million input / output tokens. Monthly billing is higher where a plan offers it: Claude Team and ChatGPT Business standard are **\$25**, and premium is **\$125**.

October 2026 list prices from Claude, OpenAI, Cursor, and GitHub. Claude Enterprise starts at 20 seats. ChatGPT Business starts at 2. ChatGPT Enterprise is a sales quote. These are not measurements from this repository.

What moves onto the GPU

CLASSIFY THE REQUEST, THEN ROUTE




Local path

- **OpenCode** edits files, runs tools, and stays on the workstation network.
- **vLLM on ROCm** serves Qwen3.8-27B, one replica per GPU.
- Prompts and repository context do not leave the building.
- After the hardware is paid for, another token does not generate another invoice.

What the dollars mean

- The 29 Sep 2026 plan in [docs/TC0.md](#). Card prices are supplied estimates.
- R9700S and MI350P rates are measured. R9600D rates are a 150 W projection.
- The local 27B model is the offload target, not a priced substitute for a frontier model.


Buy the hardware once



	8x R9700S	8x MI350P	2x R9700
Capex	\$49,000	\$209,000	\$6,500
3-year power	\$3,787	\$7,574	in total
3-year cost	\$52,787	\$216,574	\$8,235
Per year	\$17,596	\$72,191	\$2,745

- DRAM matches the GPU memory. Power is the GPU board only, at **\$0.12/kWh** and 50% of nameplate, for 26,298 hours.
- Facility power and residual value are excluded.
- Eight R9700S cards are the low capex fill. Electricity is a few percent of that bill.
- The MI350P server costs about four times as much over three years.
- The workstation is a different host: \$6,500 capex and a 550 W planning draw.

Price of a million output tokens



Shape	Sonnet 5	8x R9700S	8x MI350P
8,192 / 1,024	\$26.00	\$0.68	\$1.92
1,024 / 1,024	\$12.00	\$0.57	\$0.98
1,024 / 8,192	\$10.25	\$0.59	\$1.41

- Sonnet 5 list rate on the same input and output shape.
- On-premises prices are measured full-server throughput at concurrency 4, spread over every hour of three years.
- An 8x5 workday multiplies those on-premises prices by 4.2. The R9700S 8,192/1,024 cell becomes \$2.87, against \$26.
- At batch C8 on that shape, eight R9700S cards reach \$0.50. Production MI350P is still \$1.00.

The card pays for the seats it replaces

Break-even against the API

One R9700S share is **\$183 per month**. At the Sonnet 5 8,192/1,024 rate, that matches about **7 million** output tokens.

R9700S, 8,192/1,024	Workday month	API bill
C4, 102 tok/s	63.5 million tokens	\$1,651
Hardware share		\$183

The share is covered at about **11%** of that C4 workday. Tokens past that point are the reduction.

Seats the invoice avoids


This is list price avoided, not a claim that the local 27B matches the hosted model.

3-year hardware	Max or Ultra	Cursor Teams
2× R9700, \$8,235	1	6
8× R9700S, \$52,787	7	36
8× MI350P, \$216,574	30	150

A \$20 Pro seat is the wrong comparison. The reduction shows up on a high cap, or on API rates after the included pool is gone.

Where the larger GPU earns it

R9700S flattens once the 32 GB card protects latency. MI350P is still climbing at C32: **12,035 tok/s** and **\$0.19** per million on 1,024/1,024, or **\$0.80** on an 8×5 duty cycle.



ROUTE TO THE RIGHT MODEL. NOT THE BIGGEST.

INCOMING REQUESTS

Explain this function

Low complexity

Fix a type error

Standard coding task

Private data

Hard policy: stay local

ROUTE TO BEST FIT

Small model

Fast, if quality holds

Mid model

Local 27B, most code

Local only

Policy blocks the frontier

THE PAYOFF

38x

lower \$ / million
vs Sonnet 5 at C4

11%

of a C4 workday
covers the GPU share

Quality floor

before the cheap route

Which box to buy

Workstation: a few developers are on Max, Ultra, or API overage, and the repository has to stay inside the network. Three years of a dual R9700 is about one Max seat.

8x R9700S: a team shares one interactive coding agent and per-GPU concurrency stays near C4 to C8. This is the lowest measured dollar per million tokens in that band.

8x MI350P: the same server has to hold a larger model, or stay busy at high concurrency, where its token price keeps falling.

Sources

- On-premises capex, electricity, and token prices: `docs/TC0.md`, 29 September 2026.
- Figures: `docs/figures/tco/`. Rebuild with `python3 scripts/generate_tco_plots.py`. Routing diagrams: `docs/assets/`.
- Dual-R9700 workstation annual cost: `docs/PDD-FRAMEWORK.md`.
- Local agent path: OpenCode plus vLLM, described in the repository README. Per-request local versus frontier routing: `scripts/jev_gateway.py`.
- Routing method on slide 8: classify the request, apply hard policy and capability rules, then choose the least expensive model that still clears a quality floor. A cascade escalates a failed check. A fallback covers an outage. Method from the [LLM Routing Guide](#), 14 September 2026. The 38x and 11% figures are from this deck.
- R9600D token prices in the TCO note are scaled from R9700S measurements. RTX PRO 6000 has no priced cell in these three shapes.
- List prices, October 2026: [Claude plans](#), [Claude Team](#), [Claude Enterprise](#), [ChatGPT Business and Enterprise](#), [Cursor plans](#), [GitHub Copilot plans](#). Sonnet 5 at \$2 / \$10 per million input / output tokens is the published API rate used for the shape comparison.

```
./scripts/build_presentation.sh coder-tco
```